

Machine Learning Project 22-23

Reza Gonabadi
Zeinab Mohammadpour

- First of all, We should import most needed library and package :

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import f1_score
from sklearn.model_selection import train_test_split
# During the project we will add some necessary library
```

- Our dataset is in the CSV format and we use pandas function like read_csv to import data to our environment

```
Data = pd.read_csv(r'C:\Reza Gonabadi\Python Code\Polimi Project\assignment_part1\tyres_train.csv')
data = pd.DataFrame(Data)
```

```
#Importing The Test Set
tyres_test = pd.read_csv(r'C:\Reza Gonabadi\Python Code\Polimi Project\tyres_test.csv')
tyres_test = pd.DataFrame(tyres_test)
print("the test set size is : ", tyres_test.shape)
```

```
the test set size is : (7984, 15)
```

- The basic things we should do are knowing about size of our data

```
data.shape

(3000, 16)
```

- As you can see we have 3000 rows which shows number of sample and 16 columns which shows features and labels of classes; the next step should describe our data, what are our features and labels names and how they present in the dataset*

Exploratory Data Analysis

```
#gives information about the data types, columns, null value counts, memory usage etc
data.info(verbose=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   vulc                  3000 non-null  float64
 1   perc_nat_rubber       3000 non-null  int64  
 2   wiring_strength       3000 non-null  int64  
 3   weather               3000 non-null  float64
 4   perc_imp              3000 non-null  float64
 5   temperature           3000 non-null  float64
 6   tread_type            3000 non-null  int64  
 7   tyre_season           3000 non-null  int64  
 8   elevation             3000 non-null  float64
 9   month                 3000 non-null  int64  
10   tread_depth           3000 non-null  int64  
11   tyre_quality          3000 non-null  int64  
12   perc_exp_comp         3000 non-null  float64
13   diameter              890 non-null   float64
14   add_layers            3000 non-null  int64  
15   failure               3000 non-null  int64  
dtypes: float64(7), int64(9)
memory usage: 375.1 KB
```

```
data.isna().sum()
```

```

vulc          0
perc_nat_rubber 0
wiring_strength 0
weather       0
perc_imp      0
temperature   0
tread_type    0
tyre_season   0
elevation     0
month         0
tread_depth   0
tyre_quality  0
perc_exp_comp 0
diameter      2110
add_layers    0
failure       0
dtype: int64

```

```

#basic statistic details about the data
data.describe(include="all")

```

	vulc	perc_nat_rubber	wiring_strength	weather	perc_imp	temperature	tread_type	tyre_season	elevation	n
count	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000
mean	18.184712	31.249667	0.631333	0.282987	0.014550	-2.375360	1.618000	0.547667	376.184000	5.77
std	1.587193	4.933300	0.546673	0.183252	0.014262	5.672184	1.487553	0.497806	151.149861	3.31
min	12.312000	18.000000	0.000000	0.030000	0.000000	-19.280000	0.000000	0.000000	41.500000	0.00
25%	17.241500	28.000000	0.000000	0.160000	0.010000	-6.960000	0.000000	0.000000	263.500000	3.00
50%	17.834000	31.000000	1.000000	0.210000	0.010000	-2.080000	1.000000	1.000000	342.000000	6.00
75%	18.934000	35.000000	1.000000	0.370000	0.020000	0.080000	3.000000	1.000000	471.625000	8.00
max	29.932000	46.000000	2.000000	0.930000	0.050000	37.000000	4.000000	1.000000	939.500000	11.00

- those two function (isna().sum() , info) show us number of null features in the dataset and the last function present dataset in rows and columns, so because the number of diameter feature has too much null number it would be better to remove this feature from our dataset and work with others*

- In the section below we remove column 'diameter' from our dataset*

```

data.drop(['diameter'], axis = 1, inplace=True)
tyres_test.drop(['diameter'], axis = 1, inplace=True)

```

```

# As you can see, the diameter feature was removed
data.head(10)


```

	vulc	perc_nat_rubber	wiring_strength	weather	perc_imp	temperature	tread_type	tyre_season	elevation	month	tread_depth	tyre
0	17.990	26	1	0.16	0.01	-8.12	0	1	332.5	8	0	
1	20.704	36	1	0.30	0.01	-4.52	2	0	328.0	11	1	
2	19.156	34	1	0.30	0.01	-1.08	0	0	247.0	0	1	
3	16.802	35	1	0.19	0.02	7.44	1	0	408.0	7	3	
4	17.140	23	2	0.39	0.01	30.52	0	1	308.0	2	2	
5	20.042	38	0	0.04	0.01	-0.20	2	1	160.5	0	0	
6	21.172	33	1	0.39	0.01	-2.28	0	0	470.0	3	0	
7	16.706	32	0	0.62	0.05	-3.96	3	1	148.5	7	2	
8	17.616	25	1	0.16	0.01	-6.88	0	1	103.5	8	0	
9	17.370	34	0	0.27	0.01	-1.28	2	0	345.5	5	2	

```
tyres_test.head(10)
```

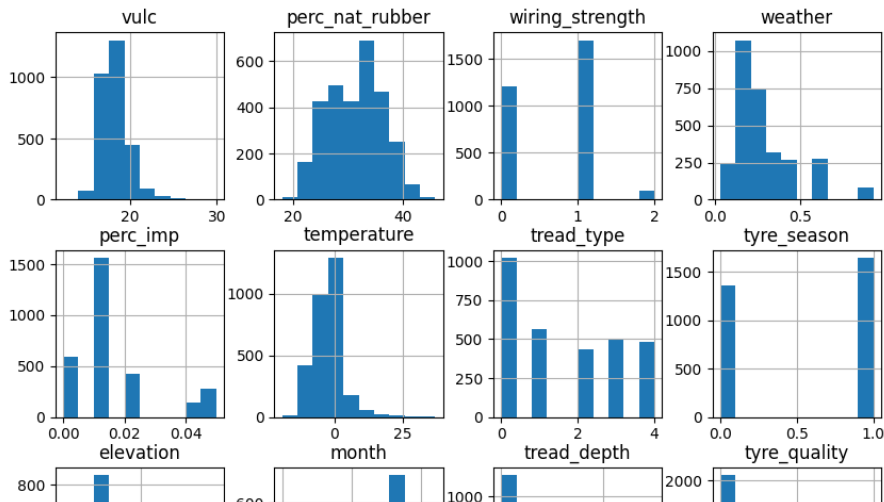
	vulc	perc_nat_rubber	wiring_strength	weather	perc_imp	temperature	tread_type	tyre_season	elevation	month	tread_depth	tyre
0	17.180	30	1	0.21	0.00	-9.24	0	1	460.5	4	3	
1	17.744	24	1	0.16	0.01	-9.12	0	1	278.5	8	0	
2	16.930	34	0	0.27	0.01	3.64	2	0	733.5	11	2	
3	22.428	34	1	0.03	0.00	0.56	3	0	235.5	9	3	
4	16.818	29	1	0.06	0.00	-0.96	3	0	461.0	9	0	
5	17.284	27	1	0.16	0.01	-11.76	4	0	251.5	7	0	
6	20.050	32	1	0.30	0.01	-4.24	0	0	552.5	0	3	
7	17.932	24	1	0.16	0.01	-7.48	0	1	283.5	8	0	
8	16.486	33	1	0.62	0.05	-1.84	2	1	305.0	6	1	

```
%matplotlib inline
data.boxplot(rot=90)
```



```
data.hist(figsize=(10,10))
```

```
array([[<AxesSubplot: title={'center': 'vulc'}>,
       <AxesSubplot: title={'center': 'perc_nat_rubber'}>,
       <AxesSubplot: title={'center': 'wiring_strength'}>,
       <AxesSubplot: title={'center': 'weather'}>],
      [ <AxesSubplot: title={'center': 'perc_imp'}>,
        <AxesSubplot: title={'center': 'temperature'}>,
        <AxesSubplot: title={'center': 'tread_type'}>,
        <AxesSubplot: title={'center': 'tyre_season'}>],
      [ <AxesSubplot: title={'center': 'elevation'}>,
        <AxesSubplot: title={'center': 'month'}>,
        <AxesSubplot: title={'center': 'tread_depth'}>,
        <AxesSubplot: title={'center': 'tyre_quality'}>],
      [ <AxesSubplot: title={'center': 'perc_exp_comp'}>,
        <AxesSubplot: title={'center': 'add_layers'}>,
        <AxesSubplot: title={'center': 'failure'}>, <AxesSubplot: >]],
      dtype=object)
```




```
#Print class freq. through pandas: we group the data by the column target and we count the number of rows
target_dist=data.groupby('failure').size()
print(target_dist)
```

```
%matplotlib inline
```

```
#Visualize Class Counts
target_dist.plot.bar(x='',y='',title='Distribution of failure')
```

```
failure
0    1992
1    1008
dtype: int64
<AxesSubplot: title={'center': 'Distribution of failure'}, xlabel='failure'>
```



- the Distribution of failure above shows us that our dataset classes are not well-balance so we will oversample the data in the next steps

▼ Data Preparation

- We want to remove outliers in dataset using standard deviations technique in this section we remove every sample which its feature greater than 3 in STD

```
def remove_outliers(df,columns,n_std):
    for col in columns:
        print('Working on column: {}'.format(col))

        mean = df[col].mean()
        sd = df[col].std()

        df = df[(df[col] <= mean+(n_std*sd))]

    return df

data = remove_outliers(data,['vulc','perc_nat_rubber','weather','temperature','perc_imp','wiring_strength','elevation','perc_exp_comp'],3)
# data.failure.value_counts()

Working on column: vulc
Working on column: perc_nat_rubber
Working on column: weather
Working on column: temperature
Working on column: perc_imp
Working on column: wiring_strength
Working on column: elevation
Working on column: perc_exp_comp
```

- As you see before our dataset was not well balanced in classes, so it would be better to oversample it first and then using it in our machines

```
from sklearn.utils import resample

#Over-sample Minority Class
#1) Separate majority and minority classes
df_majority = data[data.failure==0] #"target" is the name of the target column, change it accordingly to your dataset
df_minority = data[data.failure==1] #"target" is the name of the target column, change it accordingly to your dataset

#2) Oversample minority class
df_minority_oversampled = resample(df_minority,
                                   replace=True,
                                   n_samples=1900,    # number of samples into the minority class (change it accordingly to your dataset)
                                   random_state=123) # reproducible results

#3) Combine oversampled minority class with majority class
df_oversampled = pd.concat([df_minority_oversampled, df_majority])

#4) Display new class counts
df_oversampled.failure.value_counts() #"target" is the name of the target column, change it accordingly to your dataset

data = df_oversampled.copy()
data
```

	vulc	perc_nat_rubber	wiring_strength	weather	perc_imp	temperature	tread_type	tyre_season	elevation	month	tread_depth	t
1677	16.384	33	0	0.37	0.02	-0.48	0	0	242.0	9	0	

- We separte X, y from dataset X means our features and y means our labels

```
X = data.drop(['failure'],axis=1)
```

```
y = data.failure.values
```

- One on the important thing before classification and every machine learning problem is Normalizing data, so we normalize the data in the next step

```
4994 19.290 29 1 0.00 0.00 -1.00 2 1 107.0 0 0
```

```
# Scale data
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler().fit(X)
```

```
scaler_tyres = StandardScaler().fit(tyres_test)
```

```
# We compute the scaler
```

```
scaled_data = scaler.transform(X.astype(float))
scaled_X = pd.DataFrame(scaled_data.astype(float))
scaled_X.columns = X.columns
```


```
scaled_tyres_test_data = scaler_tyres.transform(tyres_test.astype(float))
scaled_tyres_test = pd.DataFrame(scaled_tyres_test_data.astype(float))
scaled_tyres_test.columns = X.columns
```

```
scaled_X.boxplot()
```

```
# scaled_tyres_test.boxplot()
```

```
X = scaled_X.copy()
```

```
tyres_test = scaled_tyres_test.copy()
```



```
tyres_test.shape
```

```
(7984, 14)
```

```
X
```

	vulc	perc_nat_rubber	wiring_strength	weather	perc_imp	temperature	tread_type	tyre_season	elevation	month	tread_c
0	-1.220108	0.338219	-1.114479	0.549167	0.349774	0.375608	-1.144667	-1.064686	-0.880782	0.986629	-1.13
1	-0.306383	-1.129147	-1.114479	2.140485	2.374484	2.398950	-0.470941	0.939244	-0.405299	1.282640	1.50
2	1.086240	-1.548395	0.715462	0.612820	1.699581	1.926837	0.202786	0.939244	0.476803	-1.381465	0.62
3	-0.791156	0.338219	-1.114479	0.549167	0.349774	0.451483	0.876512	-1.064686	0.981799	-1.381465	-0.25
4	1.350662	0.547843	0.715462	0.612820	1.699581	0.232288	0.876512	-1.064686	-1.510386	0.690617	-0.25
...

- As you can see above our dataset is normalize now and we can use it in our machines

- Now we separate The data into Train and Test which we could use in our machine

```
#SPLIT DATA INTO TRAIN AND TEST SET
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, #X_scaled
                                                    test_size=0.30, #by default is 75%-25%
                                                    #shuffle is set True by default,
                                                    stratify=y,
                                                    random_state= 123) #fix random seed for replicability
```

```
print(X_train.shape)
```

```
(2683, 14)
```

▼ SVM

- We tried lots of machines and after all we decided to use SVM because it is more reliable and has more score from other ones
- We use Grid Search for finding best parameters in SVM inputs and after finding best parameters we use them in our machines

```
#DEFINE YOUR CLASSIFIER and THE PARAMETERS GRID
```

```
from sklearn.svm import SVC
```

```
classifier = SVC()
```

```
parameters = {"kernel":['linear','rbf','polynomial'], "C":[0.1,1,100], "gamma":[1], "degree":[2,3,4]}
```

```
#DEFINE YOUR GRIDSEARCH
```

```
...
```

```
GS performs an exhaustive search over specified parameter values for an estimator.
```

```
GS uses a Stratified K-Folds cross-validator
```

```
(The folds are made by preserving the percentage of samples for each class.)
```

```
If refit=True the model is retrained on the whole training set with the best found params
```

```
...
```

```
from sklearn.model_selection import GridSearchCV
```

```
gs = GridSearchCV(classifier, parameters, cv=3, scoring = 'f1', verbose=50, n_jobs=-1, refit=True)
```

```
#TRAIN YOUR CLASSIFIER
```

```
gs = gs.fit(X_train, y_train)
```

```
Fitting 3 folds for each of 27 candidates, totalling 81 fits
```

```
C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py:378: FitFa
27 fits failed out of a total of 81.
```

```
The score on these train-test partitions for these parameters will be set to nan.
```

```
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

```
Below are more details about the failures:
```

```
-----
5 fits failed with the following error:
```

```
Traceback (most recent call last):
```

```
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py",
  estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\svm\_base.py", line 180, in fit
  self._validate_params()
```

```
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 570, in _validate
  validate_parameter_constraints()
```

```
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line
  raise InvalidParameterError()
```

```

sklearn.utils._param_validation.InvalidParameterError: The 'kernel' parameter of SVC must be a str among {'rbf', 'linear', 'precom
-----
5 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py",
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\svm\_base.py", line 180, in fit
    self._validate_params()
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 570, in _validate_
    validate_parameter_constraints(
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'kernel' parameter of SVC must be a str among {'poly', 'linear', 'sigmo
-----

2 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py",
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\svm\_base.py", line 180, in fit
    self._validate_params()
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 570, in _validate_
    validate_parameter_constraints(
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'kernel' parameter of SVC must be a str among {'sigmoid', 'linear', 'pr
-----

1 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py",
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\svm\_base.py", line 180, in fit
    self._validate_params()
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 570, in _validate_
    validate_parameter_constraints(
  File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line

#summarize the results of your GRIDSEARCH
print('***GRIDSEARCH RESULTS***')

print("Best score: %f using %s" % (gs.best_score_, gs.best_params_))
means = gs.cv_results_['mean_test_score']
stds = gs.cv_results_['std_test_score']
params = gs.cv_results_['params']

for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

***GRIDSEARCH RESULTS***
Best score: 0.777330 using {'C': 100, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
0.723492 (0.006207) with: {'C': 0.1, 'degree': 2, 'gamma': 1, 'kernel': 'linear'}
0.723221 (0.000870) with: {'C': 0.1, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 0.1, 'degree': 2, 'gamma': 1, 'kernel': 'polynomial'}
0.723492 (0.006207) with: {'C': 0.1, 'degree': 3, 'gamma': 1, 'kernel': 'linear'}
0.723221 (0.000870) with: {'C': 0.1, 'degree': 3, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 0.1, 'degree': 3, 'gamma': 1, 'kernel': 'polynomial'}
0.723492 (0.006207) with: {'C': 0.1, 'degree': 4, 'gamma': 1, 'kernel': 'linear'}
0.723221 (0.000870) with: {'C': 0.1, 'degree': 4, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 0.1, 'degree': 4, 'gamma': 1, 'kernel': 'polynomial'}
0.725227 (0.004027) with: {'C': 1, 'degree': 2, 'gamma': 1, 'kernel': 'linear'}
0.774696 (0.015731) with: {'C': 1, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 1, 'degree': 2, 'gamma': 1, 'kernel': 'polynomial'}
0.725227 (0.004027) with: {'C': 1, 'degree': 3, 'gamma': 1, 'kernel': 'linear'}
0.774696 (0.015731) with: {'C': 1, 'degree': 3, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 1, 'degree': 3, 'gamma': 1, 'kernel': 'polynomial'}
0.725227 (0.004027) with: {'C': 1, 'degree': 4, 'gamma': 1, 'kernel': 'linear'}
0.774696 (0.015731) with: {'C': 1, 'degree': 4, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 1, 'degree': 4, 'gamma': 1, 'kernel': 'polynomial'}
0.726991 (0.004164) with: {'C': 100, 'degree': 2, 'gamma': 1, 'kernel': 'linear'}
0.777330 (0.014885) with: {'C': 100, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 100, 'degree': 2, 'gamma': 1, 'kernel': 'polynomial'}
0.726991 (0.004164) with: {'C': 100, 'degree': 3, 'gamma': 1, 'kernel': 'linear'}
0.777330 (0.014885) with: {'C': 100, 'degree': 3, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 100, 'degree': 3, 'gamma': 1, 'kernel': 'polynomial'}
0.726991 (0.004164) with: {'C': 100, 'degree': 4, 'gamma': 1, 'kernel': 'linear'}
0.777330 (0.014885) with: {'C': 100, 'degree': 4, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 100, 'degree': 4, 'gamma': 1, 'kernel': 'polynomial'}

```



```
#TEST ON YOUR TEST SET
best_model = gs.best_estimator_
y_pred = best_model.predict(X_test)

y_pred_train = best_model.predict(X_train)

y_tyres_test = best_model.predict(tyres_test)

#EVALUATE YOUR PREDICTION (on the y_test that you left aside)
from sklearn.metrics import f1_score

print('***RESULTS ON TRAIN SET***')
print("f1_score: ", f1_score(y_train, y_pred_train))
print("--")
print('***RESULTS ON TEST SET***')
print("f1_score: ", f1_score(y_test, y_pred))

***RESULTS ON TRAIN SET***
f1_score: 1.0
--
***RESULTS ON TEST SET***
f1_score: 0.8576923076923076
```

▼ As you can see the f1 score in test set data is 85% and in train set is 100%


```
#PRINT SOME FURTHER METRICS
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.82	0.96	0.88	581
1	0.95	0.78	0.86	570
accuracy			0.87	1151
macro avg	0.88	0.87	0.87	1151
weighted avg	0.88	0.87	0.87	1151

```
#CONFUSION MATRIX
from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred))
```

```
[[557 24]
 [124 446]]
```

```
# Plot confusion matrix
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap="Blues"); #annot=True to annotate cells fmt: format
```



```

from sklearn import metrics

model = SVC(C=0.1, gamma=0.0001, kernel='linear', probability=True)

model.fit(X_train, y_train)
y_pred = model.predict(X_test)


y_probs = model.predict_proba(X_test) #predict_proba gives the probabilities for the target (0 and 1 in your case)

fpr, tpr, thresholds=metrics.roc_curve(y_test, y_probs[:,1])

import matplotlib.pyplot as plt
plt.plot(fpr, tpr, label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()

auc = metrics.roc_auc_score(y_test, y_probs[:,1])
print('AUC: %.2f' % auc)

```



AUC: 0.81

- We use all 14 features in our machine. It would be better if we apply feature selection technique to reduce size of the features. In that way our machine would learn faster and better than using all 14 features.

▼ Apply PCA


```

#PCA fit
from sklearn.decomposition import PCA
pca = PCA()
pca.fit(scaled_X)
df_pca = pd.DataFrame(pca.transform(scaled_X))


pca_tyres_test = PCA()
pca_tyres_test.fit(tyres_test)
df_pca_tyres_test = pd.DataFrame(pca_tyres_test.transform(tyres_test))

explained_variance=pd.DataFrame(pca.explained_variance_ratio_)
%matplotlib inline
import seaborn as sns
ax = sns.barplot( data=explained_variance.transpose())

```



```
explained_variance_tyres=pd.DataFrame(pca_tyres_test.explained_variance_ratio_)
%matplotlib inline
import seaborn as sns
ax = sns.barplot( data=explained_variance_tyres.transpose())
```



- Based on this plot I believe that the first 6 pca's are enough for choosing in our machine

```
pd.DataFrame(pca.components_, columns=X.columns)
```

```

vulc  perc nat rubber  wiring strength  weather  perc imp  temperature  tread type  tvre season  elevation  month  tread der
X_pca = pd.DataFrame(data = pca.transform(X)
                    ,columns = ['pc1', 'pc2','pc3','pc4','pc5', 'pc6','pc7','pc8', 'pc9','pc10','pc11','pc12', 'pc13','pc14'])

X_train_pca = pd.DataFrame(data = pca.transform(X_train)
                    ,columns = ['pc1', 'pc2','pc3','pc4','pc5', 'pc6','pc7','pc8', 'pc9','pc10','pc11','pc12', 'pc13','pc14'])

X_test_pca = pd.DataFrame(data = pca.transform(X_test)
                    ,columns = ['pc1', 'pc2','pc3','pc4','pc5', 'pc6','pc7','pc8', 'pc9','pc10','pc11','pc12', 'pc13','pc14'])

tyres_test_pca = pd.DataFrame(data = pca_tyres_test.transform(tyres_test)
                    ,columns = ['pc1', 'pc2','pc3','pc4','pc5', 'pc6','pc7','pc8', 'pc9','pc10','pc11','pc12', 'pc13','pc14'])


-----

import matplotlib.pyplot as plt

x = X_train_pca.iloc[:,0]
y = X_train_pca.iloc[:,1]

plt.scatter(x, y,alpha=0.2,c=y_train )
plt.show()

```




```

import matplotlib.pyplot as plt

x = X_test_pca.iloc[:,0]
y = X_test_pca.iloc[:,1]

plt.scatter(x, y,alpha=0.2,c=y_test )
plt.show()

```



▼ Using 6 PC's for the Machines

```
#DEFINE YOUR CLASSIFIER and THE PARAMETERS GRID
from sklearn.svm import SVC
```

```
classifier = SVC()
parameters = {'kernel':['linear','rbf','polynomial'], "C":[0.1,1,100],"gamma":[1], "degree":[2,3,4]}
```

```
#DEFINE YOUR GRIDSEARCH
'''
```

```
GS perfoms an exhaustive search over specified parameter values for an estimator.
GS uses a Stratified K-Folds cross-validator
(The folds are made by preserving the percentage of samples for each class.)
If refit=True the model is retrained on the whole training set with the best found params
'''
```

```
from sklearn.model_selection import GridSearchCV
gs = GridSearchCV(classifier, parameters, cv=3, scoring = 'f1', verbose=50, n_jobs=-1, refit=True)
```

```
#TRAIN YOUR CLASSIFIER
gs = gs.fit(X_train_pca.iloc[:, :6], y_train)
```

```
Fitting 3 folds for each of 27 candidates, totalling 81 fits
C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py:378: FitFa
27 fits failed out of a total of 81.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

2 fits failed with the following error:

Traceback (most recent call last):

```
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py",
  estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\svm\_base.py", line 180, in fit
  self._validate_params()
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 570, in _validate
  validate_parameter_constraints(
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line
  raise InvalidParameterError(
```

sklearn.utils._param_validation.InvalidParameterError: The 'kernel' parameter of SVC must be a str among {'sigmoid', 'rbf', 'preco

5 fits failed with the following error:

Traceback (most recent call last):

```
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py",
  estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\svm\_base.py", line 180, in fit
  self._validate_params()
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 570, in _validate
  validate_parameter_constraints(
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line
  raise InvalidParameterError(
```

sklearn.utils._param_validation.InvalidParameterError: The 'kernel' parameter of SVC must be a str among {'sigmoid', 'linear', 'pr

1 fits failed with the following error:

Traceback (most recent call last):

```
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py",
  estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\svm\_base.py", line 180, in fit
  self._validate_params()
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 570, in _validate
  validate_parameter_constraints(
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line
  raise InvalidParameterError(
```

sklearn.utils._param_validation.InvalidParameterError: The 'kernel' parameter of SVC must be a str among {'precomputed', 'rbf', 's

5 fits failed with the following error:

Traceback (most recent call last):

```
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection\_validation.py",
  estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\svm\_base.py", line 180, in fit
  self._validate_params()
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py", line 570, in _validate
  validate_parameter_constraints(
File "C:\Users\Reza.Gonabadi\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\_param_validation.py", line
```

```
#summarize the results of your GRIDSEARCH
print('***GRIDSEARCH RESULTS***')

print("Best score: %f using %s" % (gs.best_score_, gs.best_params_))
means = gs.cv_results_['mean_test_score']
stds = gs.cv_results_['std_test_score']
params = gs.cv_results_['params']

for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

***GRIDSEARCH RESULTS***
Best score: 0.778840 using {'C': 100, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
0.683551 (0.005481) with: {'C': 0.1, 'degree': 2, 'gamma': 1, 'kernel': 'linear'}
0.723810 (0.002364) with: {'C': 0.1, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 0.1, 'degree': 2, 'gamma': 1, 'kernel': 'polinomial'}
0.683551 (0.005481) with: {'C': 0.1, 'degree': 3, 'gamma': 1, 'kernel': 'linear'}
0.723810 (0.002364) with: {'C': 0.1, 'degree': 3, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 0.1, 'degree': 3, 'gamma': 1, 'kernel': 'polinomial'}
0.683551 (0.005481) with: {'C': 0.1, 'degree': 4, 'gamma': 1, 'kernel': 'linear'}
0.723810 (0.002364) with: {'C': 0.1, 'degree': 4, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 0.1, 'degree': 4, 'gamma': 1, 'kernel': 'polinomial'}
0.685512 (0.006622) with: {'C': 1, 'degree': 2, 'gamma': 1, 'kernel': 'linear'}
0.778506 (0.017438) with: {'C': 1, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 1, 'degree': 2, 'gamma': 1, 'kernel': 'polinomial'}
0.685512 (0.006622) with: {'C': 1, 'degree': 3, 'gamma': 1, 'kernel': 'linear'}
0.778506 (0.017438) with: {'C': 1, 'degree': 3, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 1, 'degree': 3, 'gamma': 1, 'kernel': 'polinomial'}
0.685512 (0.006622) with: {'C': 1, 'degree': 4, 'gamma': 1, 'kernel': 'linear'}
0.778506 (0.017438) with: {'C': 1, 'degree': 4, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 1, 'degree': 4, 'gamma': 1, 'kernel': 'polinomial'}
0.685024 (0.007205) with: {'C': 100, 'degree': 2, 'gamma': 1, 'kernel': 'linear'}
0.778840 (0.012534) with: {'C': 100, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 100, 'degree': 2, 'gamma': 1, 'kernel': 'polinomial'}
0.685024 (0.007205) with: {'C': 100, 'degree': 3, 'gamma': 1, 'kernel': 'linear'}
0.778840 (0.012534) with: {'C': 100, 'degree': 3, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 100, 'degree': 3, 'gamma': 1, 'kernel': 'polinomial'}
0.685024 (0.007205) with: {'C': 100, 'degree': 4, 'gamma': 1, 'kernel': 'linear'}
0.778840 (0.012534) with: {'C': 100, 'degree': 4, 'gamma': 1, 'kernel': 'rbf'}
nan (nan) with: {'C': 100, 'degree': 4, 'gamma': 1, 'kernel': 'polinomial'}

#TEST ON YOUR TEST SET
best_model = gs.best_estimator_
y_pred = best_model.predict(X_test_pca.iloc[:,6])

y_pred_train = best_model.predict(X_train_pca.iloc[:,6])

y_tyres_test_result = best_model.predict(tyres_test_pca.iloc[:,6])

#EVALUATE YOUR PREDICTION (on the y_test that you left aside)
from sklearn.metrics import f1_score

print('***RESULTS ON TRAIN SET***')
print("f1_score: ", f1_score(y_train, y_pred_train))
print("--")
print('***RESULTS ON TEST SET***')
print("f1_score: ", f1_score(y_test, y_pred))

***RESULTS ON TRAIN SET***
f1_score: 1.0
--
***RESULTS ON TEST SET***
f1_score: 0.8266666666666668
```

As you can see the f1 score in the test set data is 82.7% and in the train set is 100%, and I consider this as a final score of this project.

```
#PRINT SOME FURTHER METRICS
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))

          precision    recall  f1-score   support

0               0.86         0.77         0.81         581
```

	1	0.79	0.87	0.83	570
accuracy				0.82	1151
macro avg		0.82	0.82	0.82	1151
weighted avg		0.82	0.82	0.82	1151


```
#CONFUSION MATRIX
```

```
from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred))
```

```
[[447 134]
 [ 74 496]]
```

```
# Plot confusion matrix
```

```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap="Blues"); #annot=True to annotate cells fmt: format
```



```
count = (y_tyres_test_result == 1).sum()
count
```

```
2529
```

```
arr = y_tyres_test_result
int_array = arr.astype('int')

np.savetxt("result.txt", int_array)
```